

ECON 706: PROBLEM SET 3

DANIEL PFEFFER

Problem 1. Write a program which takes as inputs 2 uniformly distributed random variables and uses the Box-Muller approach to obtain 2 standard normally distributed random variables.

Solution. The function `box_muller` expects two uniform random variables of arbitrary size and generates realizations of two standard normal random variables using the Box-Muller approach.

```
import numpy as np
import matplotlib.pyplot as plt

def box_muller(u1, u2):
    """u1, u2 are uniform random variables of size T
    Returns two N(0,1) random variables of size T"""
    z1 = np.sqrt(-2*np.log(1-u1))*np.cos(2*np.pi*u2)
    z2 = np.sqrt(-2*np.log(1-u1))*np.sin(2*np.pi*u2)
    return np.array([z1, z2])
```

Problem 2. Consider the following 2-variable VAR(1):

$$y_t = \Phi y_{t-1} + \epsilon_t, \quad \epsilon_t \sim N(0, \Sigma),$$

where $y_t = (y_{1t}, y_{2t})'$, $\epsilon_t = (\epsilon_{1t}, \epsilon_{2t})'$,

$$\Phi = \begin{pmatrix} 0.6 & 0 \\ 0.9 & 0.3 \end{pmatrix},$$

and where Σ is such that $\text{var}(y_{1t}) = 2$, $\text{var}(y_{2t}) = 8$ and $\text{corr}(y_{1t}, y_{2t}) = 0.5$.

For $T = 1000$, use your program from exercise 1 to simulate $\{y_{1t}, y_{2t}\}_{t=1}^T$. Draw y_0 from the unconditional distribution of y_t .

Solution. The covariance matrix is

$$\Sigma = \begin{pmatrix} 2 & 2 \\ 2 & 8 \end{pmatrix},$$

since $\text{cov}(y_{1t}, y_{2t}) = \sqrt{2 \cdot 8} \cdot .5 = 2$.¹

We then generate two $\text{unif}([0, 1])$ random variables of size $T = 1000$ and call `box_muller` to simulate $x \sim N(0, I)$ where $x = (x_1, x_2)'$. Then note that since every covariance matrix Σ is positive-semidefinite, there exists a Cholesky factorization $\Sigma = PP'$, where P is a lower triangular matrix. So

$$\epsilon_t = Px \sim N(0, \Sigma).$$

Date: May 27, 2020.

¹This result is obtained directly from the fact that

$$\text{corr}(y_{1t}, y_{2t}) = \frac{\text{cov}(y_{1t}, y_{2t})}{\sigma_{y_{1t}} \sigma_{y_{2t}}},$$

Note that this also gives the initial draw for y_0 from the unconditional distribution of y_t . Finally, applying the autoregressive operator Φ to innovations ϵ_t along with the initial condition given by y_0 generates samples from the desired system. The following code implements this VAR(1) simulation.

```
# Initialize VAR(1) system parameters
Phi = [[.6, 0], [.3, .9]]
Sigma= [[2, 2], [2, 8]]

# Generate two U([0,1]) random variables
np.random.seed(42)
T = 1000
u1, u2 = np.random.rand(T), np.random.rand(T)

# Simulate N(0,I) random vector
iid_innov = box_muller(u1, u2)

# Generate innovations via Cholesky factorization
L = np.linalg.cholesky(Sigma)
eps = np.dot(L, iid_innov)
eps1, eps2 = eps[0], eps[1]

# Allocate memory for y1 and y2 series
y1, y2 = np.empty_like(eps1), np.empty_like(eps2)

# Draw y0 from unconditional distribution
y1[0], y2[0] = eps1[0], eps2[0]

# Simulate y1 and y2 series
for t in range(1, T):
    y1[t] = Phi[0][0]*y1[t-1] + eps1[t]
    y2[t] = Phi[1][0]*y1[t-1] + Phi[1][1]*y2[t-1] + eps2[t]
```

Problem 3. Estimate the univariate spectral density function of y_{1t} using the following approaches.

- Take $\hat{f}(\omega) = (2\pi)^{-1} \sum_{\tau=-T}^{T-1} \hat{\gamma}(\tau) e^{-i\omega\tau}$, where $\hat{\gamma}$ are the sample autocovariances.
- Use a uniform lag window with truncation lag of 10.
- Fit an AR(p) model for y_{1t} using a simple regression, where p is chosen based on model selection criteria, and compute the spectral density of your estimated model.

Plot and comment on the resulting estimated spectral densities.

Solution. The sample spectral density computed by first estimating the periodogram

$$I(\omega) = \frac{1}{T} \left| \sum_{t=0}^{T-1} y_t e^{2\pi i(\omega/T)t} \right|^2,$$

and then scaling the resulting values by 4π to obtain

$$\hat{f}(\omega) = \frac{1}{4\pi} I(\omega).$$

Note that this is equivalent to computing the sample autocovariance and taking its Fourier transform. This estimation procedure is implemented in the function `sample_sdf`.

```

from numpy.fft import fft

def sample_sdf(y):
    """y is an array of time domain data
    Returns w, the frequencies and f_y, the sample spectral density"""
    T = len(y)
    f_y = np.abs(fft(y))**2/T
    w = 2*np.pi*np.arange(T)/T
    w, f_y = w[:int(T/2)+1], f_y[:int(T/2)+1]
    return w, f_y/(4*np.pi)

```

The function `unif_window` estimates the spectral density function by applying a uniform window with a truncation lag of 10 to the sample spectral density.

```

def unif_window(f, trunc_lag=10):
    """f is the sample spectral density
    Return the uniform lag window sample spectral density"""
    j = int(trunc_lag/2)
    fb = f[:j]
    ft = f[-j:]
    s = np.concatenate((fb[::-1], f, ft[::-1]))
    window = np.ones(trunc_lag)
    return np.convolve(window/window.sum(), s, mode='valid')

```

The function `ar_sdf` fits an $AR(p)$ model is fit where the order p is determined by the Bayesian information criterion. The oracle property associated with this scoring method motivates this choice, since we know the law of motion used to generate the simulated y_{1t} process is of the autoregressive order 1, which is the first model within the space containing all possible combinations of possibly infinite order autoregressive models. So a local search is performed over the restricted space of autoregressive models with a maximum lag of 15. This enables us to avoid the exponential complexity associated with a global search. The autoregressive spectral density is then estimated using

$$\hat{f}(\omega) = \frac{\hat{\sigma}^2}{2\pi} \left| \frac{1}{\phi(e^{i\omega})} \right|^2$$

where ϕ is the autoregressive polynomial determined by the BIC, and $\hat{\sigma}^2$ is the sample variance.

```

from statsmodels.tsa.ar_model import ar_select_order
from scipy.signal import freqz

def ar_sdf(y, max_lag=15, i_c='bic', cst='n'):
    """y is an array of time domain data
    Returns w, the frequencies and ar_sdf, the ar sample spectral density"""
    sel = ar_select_order(y, maxlag=max_lag, ic=i_c, trend=cst)
    res = sel.model.fit()
    ar_poly = np.insert(-res.params, 0, 1)
    w, b_fft = freqz(b=1, a=ar_poly, worN=int(len(y)/2)+1)
    ar_sdf = res.sigma2*np.abs(b_fft)**2/(4*np.pi)
    return w, ar_sdf

```

We make function calls on the simulated y_{1t} process to estimate its spectral density according to these three alternative methods.

```
# Sample SDF for y1
w_y1, f_y1 = sample_sdf(y1)

# Uniform lag window SDF for y1
w_y1_unif, f_y1_unif = unif_window(f_y1)

# Autoregressive SDF estimation for y1
ar_freq_y1, b_fft_y1 = ar_sdf(y1)
```

Figure ?? shows the plots of the resulting nonparametric sample spectral density estimates. The top graph shows the sample spectral density. Although it vaguely captures what Granger called the “typical spectral shape” associated with an AR(1) with $\phi > 0$, is a substandard estimate. In other words, although appears somewhat unbiased, the resulting figure displays clear evidence of inconsistency that can be seen in the highly fluctuating values on the y -axis. The second graph shows the uniform window with a truncation lag of 10 sample spectral density. This estimate smoothed out much of this inconsistency, as can be seen by noting the change in the values on the y -axis. The third plot shows the autoregressive spectral density estimate. The BIC made the correct model specification, as expected and the conditional maximum likelihood parameter estimate is .62, which differs from the “true” parameter by only .02. For the y_{1t} simulated system, this method generates the best estimate.

Problem 4. Consider

$$x_t = y_{1t} - y_{1,t-1}$$

$$z_t = \frac{1}{3}(y_{1,t-1} + y_{1t} + y_{1,t+1}).$$

Repeat exercise 3 for x_t and z_t . Compare the spectral densities of x_t , z_t , and y_{1t} .

Solution. The following code applies the first difference filter $b(L) = 1 - L$ to y_{1t} to generate x_t and applies the moving average filter $h(L) = (L + 1 + L^{-1})/3$ to y_{1t} to generate z_t .

```
# First difference filter
xt = np.diff(y1)

# Moving average (lag 3) filter
zt = np.convolve(np.ones(3)/3, y1, mode='valid')
```

The following code shows the function calls to `sample_sdf`, `unif_window`, and `ar_sdf` used to generate estimates of the spectral densities of the x_t and z_t processes.

```
# Sample SDFs
w_xt, f_xt = sample_sdf(xt)
w_zt, f_zt = sample_sdf(zt)

# Uniform lag window SDFs
w_xt_unif, f_xt_unif = unif_window(xt)
w_zt_unif, f_zt_unif = unif_window(zt)

# Sample SDFs
w_xt, f_xt = sample_sdf(xt)
w_zt, f_zt = sample_sdf(zt)

# Uniform lag window SDFs
```




FIGURE 1

```
f_xt_unif = unif_window(f_xt)
f_zt_unif = unif_window(f_zt)

# Autoregressive SDF estimation
xt_freq, ar_sdf_xt = ar_sdf(xt)
zt_freq, ar_sdf_zt = ar_sdf(zt)
```

Figure ?? shows estimated spectral density functions for the x_t and z_t processes. Consider first the x_t process. Since the spectrum of x_t is related to the spectrum of y_t by

$$f_x(\omega) = b(e^{-i\omega})b(e^{i\omega})f_y(\omega),$$

we analyze the function

$$b(e^{-i\omega})b(e^{i\omega}) = 2 - 2\cos\omega,$$

In the space $L_2[0, \pi]$, this function attains a maximum of 4 at $\omega = \pi$ and a minimum of 0 at $\omega = 0$, and so this filter has the effect of emphasizing the high-frequency components of y_{1t} and removing components at low frequencies. But from Problem 3 we know that y_{1t} is primarily composed of low frequency components as confirmed by all three sample spectral density functions in Figure ?. It follows that all estimated sample spectral density function of x_t are all mostly flat, increasing in value with ω . As before, the estimates become smoother as the consistency of the estimators increases, which can be seen by observing the top, middle, then bottom graphs.

Consider now the z_t process. Similarly, to understand the affect of this filter on the y_{1t} process, we analyze the function

$$h(e^{-i\omega})h(e^{i\omega}) = \frac{1}{3}(1 + e^{-i\omega} + e^{i\omega})\frac{1}{3}(1 + e^{i\omega} + e^{-i\omega}) = \frac{1}{9}(2\cos\omega + 1)^2.$$

In the space $L_2[0, \pi]$, this function attains a maximum of 1 at frequency $\omega = 0$, monotonically decreasing to a minimum of 0 at frequency $\omega = 2\pi/3$. This filter therefore accentuates components at frequency 0 and removes components at frequency $2\pi/3$. The effect of including the $y_{1,t+1}$ term leads the z_t process. This effect is most clearly observed by comparing the autoregressive spectral density estimates. It should be noted however that it is difficult to decipher this impact by comparing only the sample spectral density estimate and uniform window spectral density estimate in the first and second graphs in Figure ? and Figure ?, respectively.

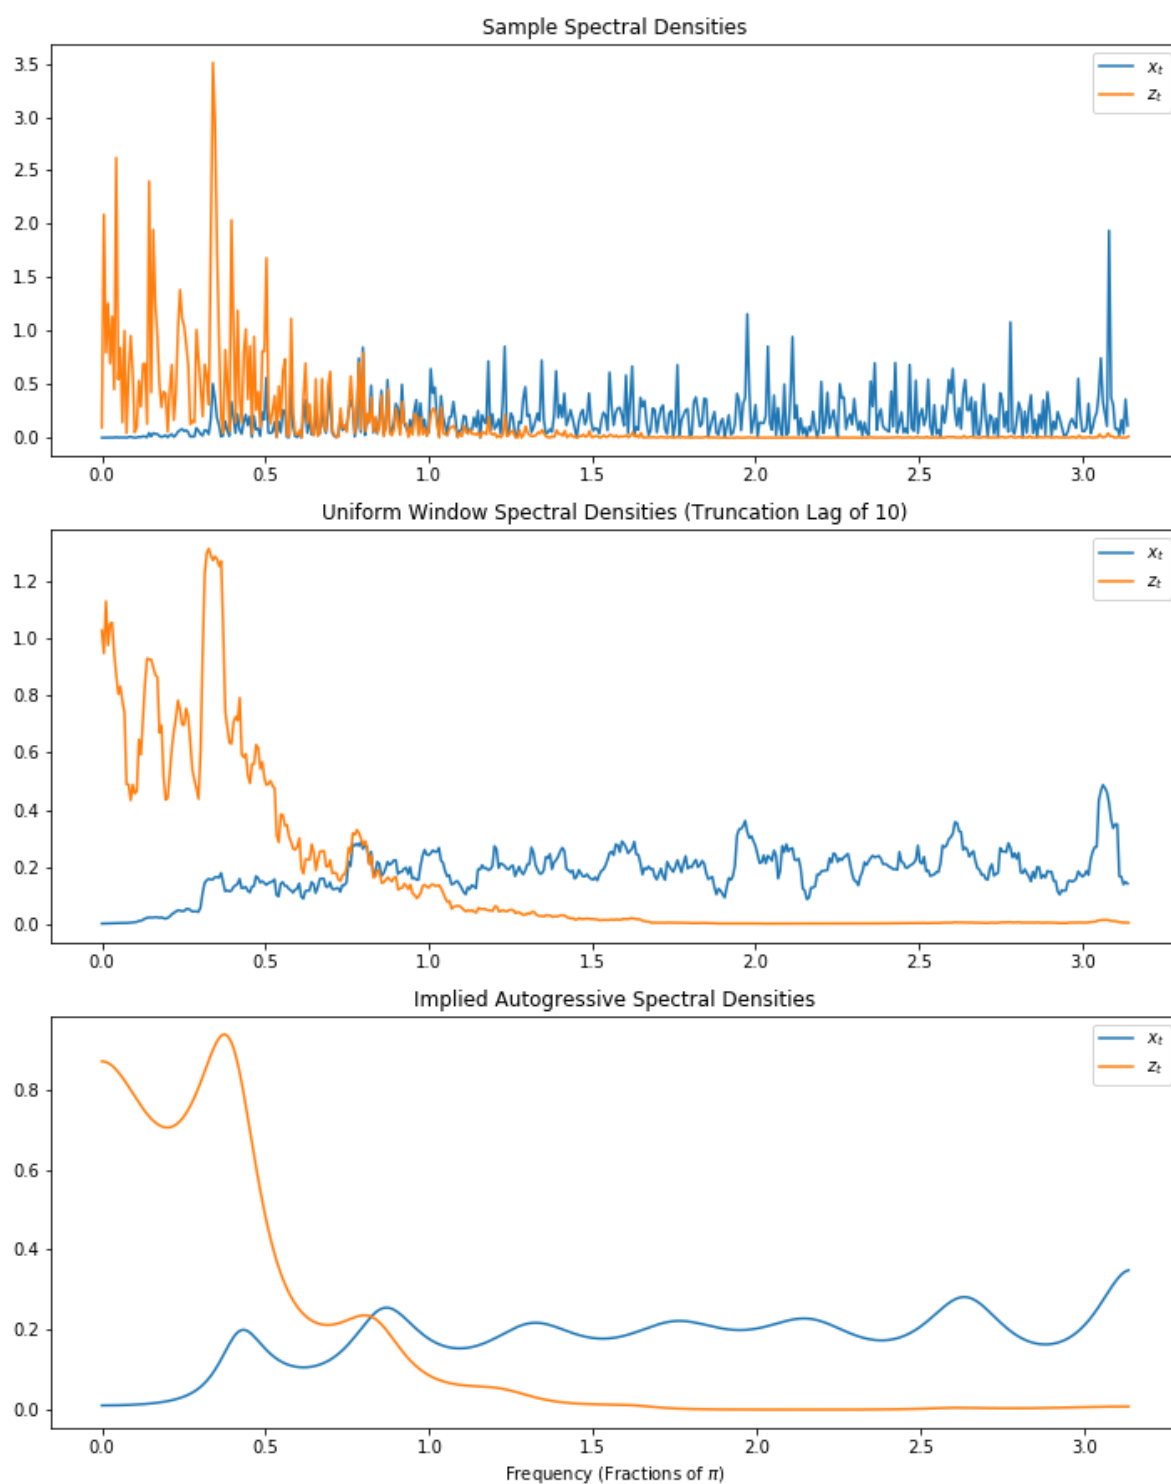


FIGURE 2